



Mobile App Ads Click Fraud Detection

Zili Chen, Chen Liang, Jinhao Jiang,
Bowen Yang and Jingyang Sui

Summary

Our project solves the problem of detecting **fraud clicks** on mobile app advertisements which induce clicks but **do not end up installing apps** and visualizes the top suspicious fraud clicks.

Motivation

The project aims to serve the app developers or companies who pay high advertisement fee for fraudulent clicks.


Approaches

Data Preprocessing

- Remove incorrect and abnormal values.
- Data unification.
- Convert data type.

Data Analysis

- Conversion rate doesn't have strong relationship with the IP frequencies.
- Fraud clicks mainly generated by certain small range of app, device, os and channels. Here we display the **CDF** of OS as an example:



Gradient Boosting Decision Trees (GBDT)

We used GBDT to train model with sigmoid objective, which built an ensemble of decision trees stage-wisely.

Innovations

- Using negative down-sampling method to deal with imbalanced datasets and improving testing AUC from 0.84 to 0.93.

- Using XGboost built-in functions to explore deeper tree branches and reduce overfitting by applying regularizations

Visualization

Submit features of the click and our model will predict whether it is a fraud click. Top suspicious fraud clicks for each feature.

Data

Dataset downloaded from Kaggle. There are 100,000 rows (observations) by 8 columns (features).

Experiments & Results

The evaluation metrics is AUC, which measures the two-dimensional area underneath the entire Receiver Operating Characteristic (ROC) curve. ROC is a graph showing **True Positive Rate (TPR)** vs. **False Positive Rate (FPR)** of a classification. We used GBDT to obtain an AUC of 0.94 from training dataset and an average AUC of 0.93 from 5 different test sets.

$$TPR = \frac{TP}{TP + FN}$$

$$FPR = \frac{FP}{FP + TN}$$

Compared with random forest model, the training AUC and testing AUC respectively decrease to 0.89 and 0.78.

